

BIOE 498/598 PJ
Experiment Design and Optimization
Spring 2021 (revised 1/24/2021)
<http://bioe498.github.io>

Course Instructor

Paul A. Jensen
pjens@illinois.edu
Office Hours: Th, 8:30-9:20am, online

Description

All engineers perform experiments. Whether they be “wet-lab” or simulated, experiments test the limits of our hypotheses and drive our understanding. Often engineers want to go beyond validating their theories and models. Engineers want the best designs – the optimal combination of cost, reliability, performance, and usability.

Part I of this course presents *Design of Experiments* and *Response Surface Methodology*, two systematic frameworks for optimizing and understanding real-world systems. Using both experimental data and computer simulations, students will explore methods to efficiently search large design spaces. Upon completion, students will be able to

- Design experiments with maximum statistical power.
- Use sequential experiments to find optimal experimental conditions.
- Tune “black box” models with multiple inputs and parameters.

Part II introduces *Reinforcement Learning*, a branch of artificial intelligence where agents both design experiment and interpret the results. Reinforcement learning underlies many recent AI breakthroughs in gameplay, logistics, protein folding, self-driving cars, and robotics. Students will learn to build agents that rival human experts when solving multistage decision problems.

The techniques in this course apply to any field of science and engineering. The course will focus on biomedical and bioengineering topics, including

- Bioprocess and metabolic engineering
- Medical device design and testing
- Optimizing molecular biology assays
- Design of synthetic gene circuits

Students will use real-world data and simulation software to optimize designs. As is typical in engineering, many of the projects lack a single “best” answer. Students will learn to balance multiple objectives and constraints when solving problems.

Audience

The course targets upper division undergraduates and graduate students in all engineering fields. Students in BIOE, ABE, and CHBE may connect best with the application areas. Examples will be selected from all undergraduate BIOE tracks. Students from MCB, ACES, and other colleges are welcome provided they meet the course prerequisites.

Prerequisites

- BIOE 210 (Linear Algebra for Biomedical Data Science)
- Graduate students with a background in linear algebra and basic statistics are welcome. These students should review Part I of the following text for background on linear statistical modeling: *Linear Algebra: Foundations of Machine Learning* by P.A. Jensen (available for free online).
- Basic computer programming skills. This course will introduce students to the R programming language – no experience with R is necessary.

Assessments

Three case studies. Students are encouraged to work in teams of two or three. Each case study is an assignment from the manager at a fictional company. Teams will receive a budget that can be spent on experiments. The deliverable for each project will be a slide deck detailing the optimization performed and team's recommendations.

Two challenge projects. In Part I, teams will use limited experiments to aim a trebuchet and hit a target. In Part II, teams will design an AI system for a computer board game.

Two in-class exams. Exams measure conceptual understanding of experimental design and optimization techniques. Any non-electronic materials are allowed during the exam, including the course textbook and notes.

Grading

Case studies: 45% ($3 \times 15\%$ each)

Challenge Projects: 30% ($2 \times 15\%$ each)

Exams: 25% ($2 \times 12.5\%$ each)

Textbooks

Design and Analysis of Experiments with R
by John Lawson; CRC Press. (required)

Regression and Other Stories
by Andrew Gelman, Jennifer Hill, & Aki Vehtari; Cambridge. (optional)

Reinforcement Learning: An Introduction (2nd edition)
by Richard Sutton & Andrew Barto; MIT Press (optional, available for free online)