GRIDWORLD + The Value function

BIOE 498/598 PJ

Markov Decision Process (MDP)

states	s	where we are
actions	a	what we choose (next state)
rewards	r	gain/loss from (s,a)

Markov Property: a state tells us everything we need to know (memoryless)

$$s_0, a_0, r_1, s_1, a_1, r_2, s_2, \cdots\cdots, s_T$$

Making Decisions in an MDP

Assume: MDP is deterministic
Process terminates (@ time T)

Let's say we are at state s_0.

$$\max \text{Reward}$$

$$= \max_a \left\{ r_1 + \sum_{i=2}^{T} r_i \right\} = \max_a \left\{ r_1 + V(s') \right\}$$

r_1: immediate reward

$\sum_{i=2}^{T} r_i$: future reward

$V(s')$: value function

GRIDWORLD

Compue V(s) for every square

If I know V(s), then I can always find the shortest path.

Get from start to finish in the fewest number of steps.

How do I compute $V(s)$?

$r_i = -1$ except at the finish.

States $S = \{1, 2, \ldots, 16\}$

start: 1 terminal: 16

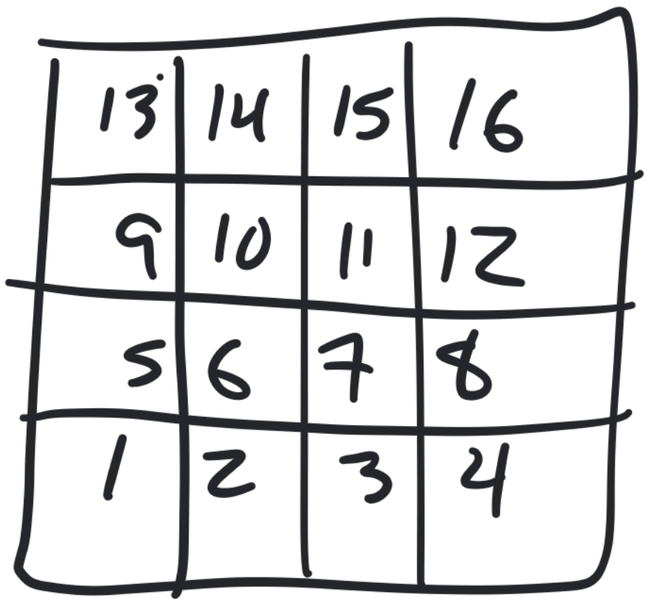

actions are defined for every state

Approach: Monte Carlo Method

1. Pick a state.
2. Take a random walk until I reach the finish
3. Count steps.
4. Repeat.

5, 6, 10, 9, 5, 6, 7, 11, 15, 16

$V(5) = -9$

V(s)

Policy Improvement

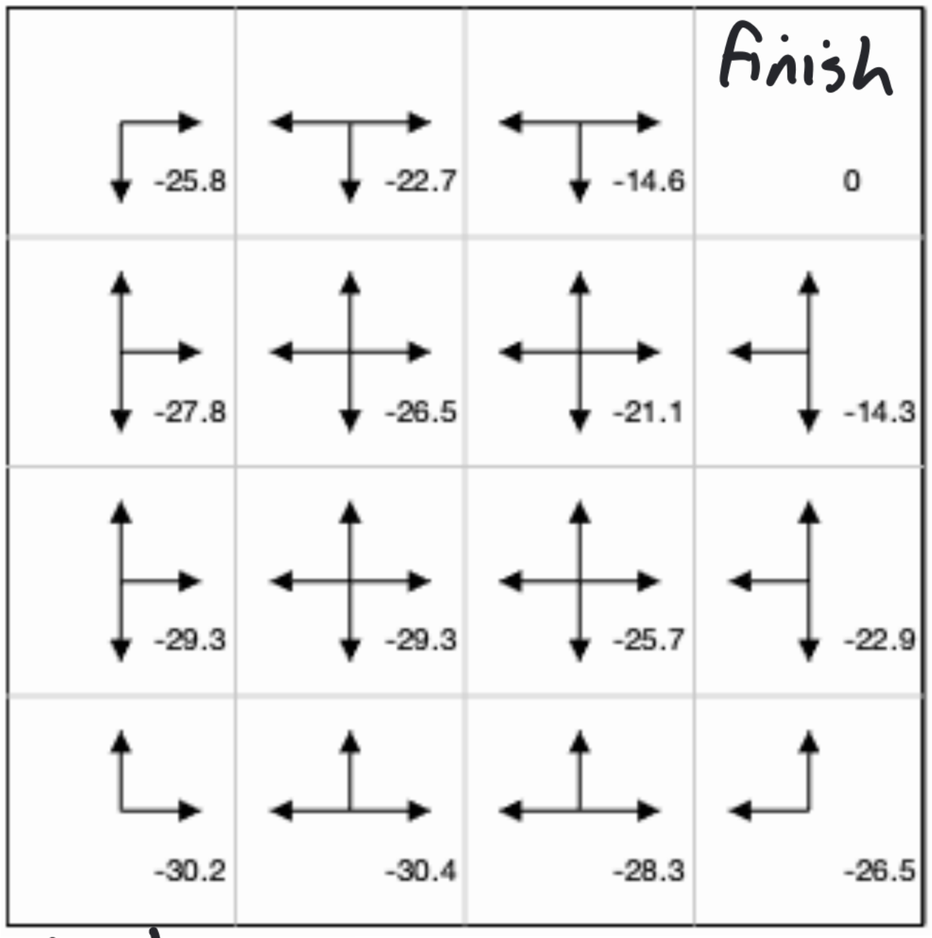

start

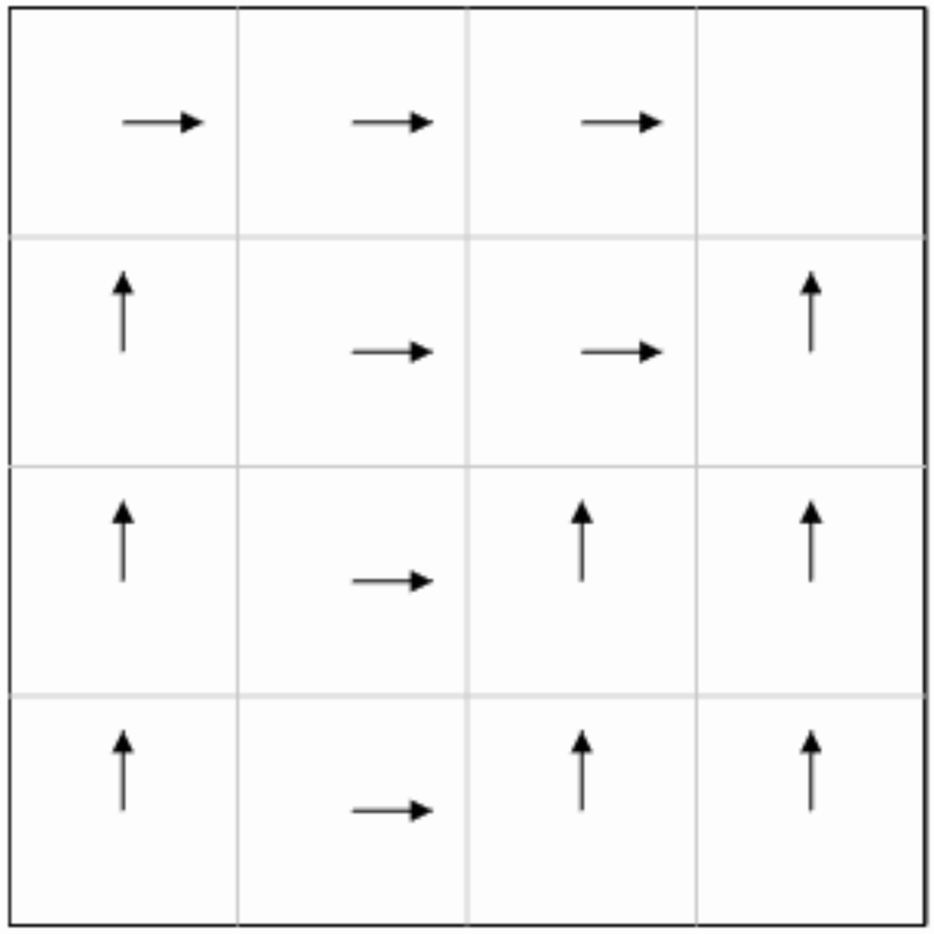

V(s)

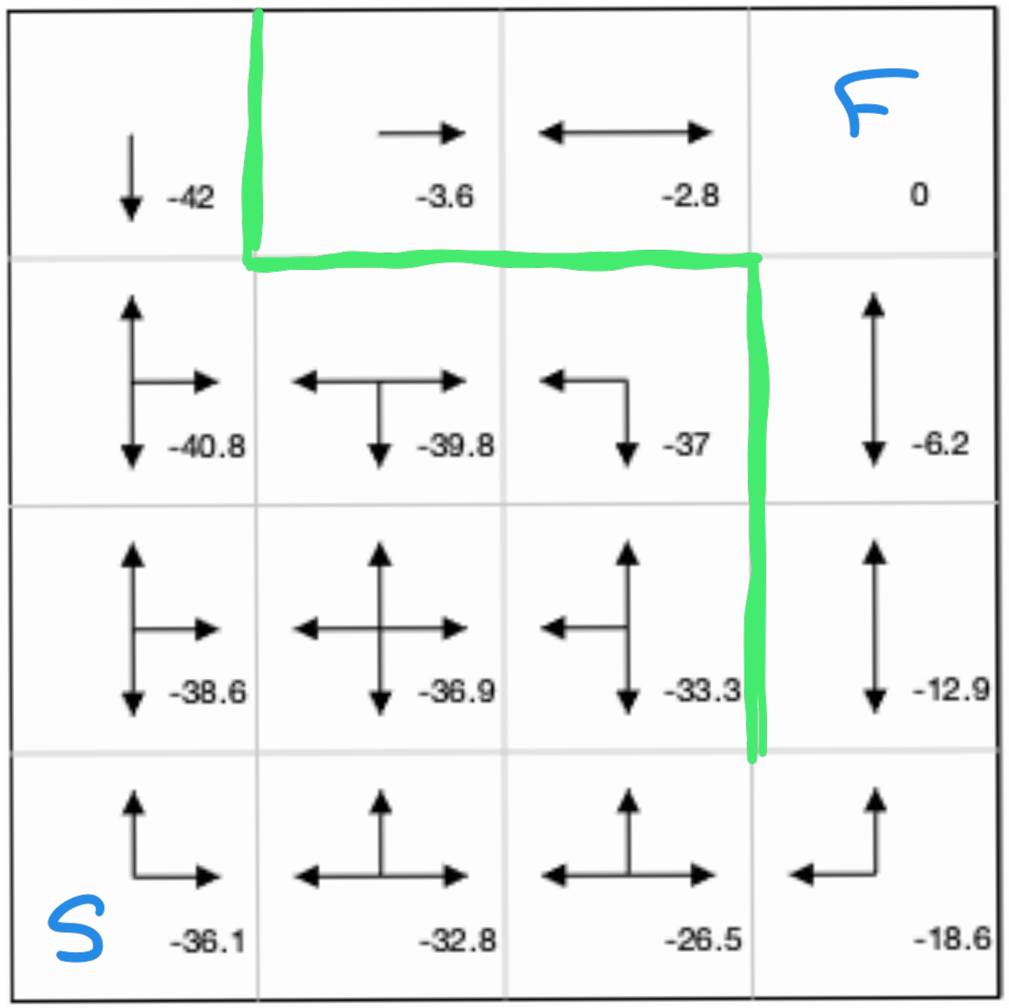

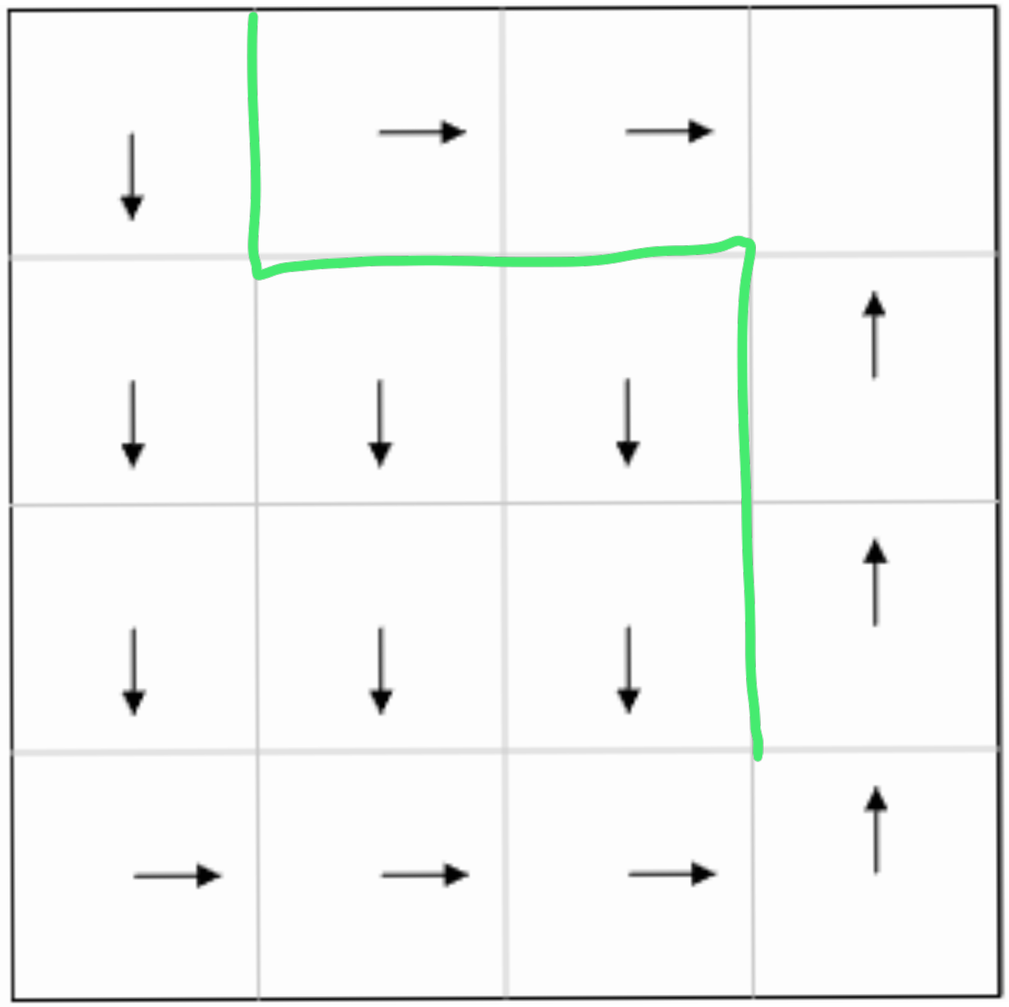

V(s)