

# Reinforcement Learning: $Q$ -learning and Tic-Tac-Go

BIOE 498/598 PJ

Spring 2021

## Review

- ▶ Discount factors shorten the horizon of RL problems, causing the agent to focus on rewards in the near future.
- ▶ Temporal Difference (TD) learning incrementally updates value functions using a new experience.
- ▶ Learning  $Q$ -factors eliminates the need to predict the next state given an action; however, the number of  $Q$ -factors is much greater than the number of states.

## Review

- ▶ Discount factors shorten the horizon of RL problems, causing the agent to focus on rewards in the near future.
- ▶ Temporal Difference (TD) learning incrementally updates value functions using a new experience.
- ▶ Learning  $Q$ -factors eliminates the need to predict the next state given an action; however, the number of  $Q$ -factors is much greater than the number of states.
- ▶ **Today:**
  - ▶ Review SARSA
  - ▶  $Q$ -learning
  - ▶ Tic-Tac-Go

## Learning $Q$ -factors

Using  $Q$ -factors, the policy problem at state  $s_i$

$$\max_a \mathbb{E} \{r_i + \gamma V(s_{i+1})\}$$

becomes

$$\max_a \mathbb{E} \{Q(s_i, a)\} .$$

## Learning $Q$ -factors

Using  $Q$ -factors, the policy problem at state  $s_i$

$$\max_a \mathbb{E} \{r_i + \gamma V(s_{i+1})\}$$

becomes

$$\max_a \mathbb{E} \{Q(s_i, a)\} .$$

- ▶ **Pro:** We do not need a model or a way to predict  $s_{i+1}$ .
- ▶ **Con:** We need to learn a  $Q$ -factor for every state/action pair.

## Learning $Q$ -factors

Using  $Q$ -factors, the policy problem at state  $s_i$

$$\max_a \mathbb{E} \{r_i + \gamma V(s_{i+1})\}$$

becomes

$$\max_a \mathbb{E} \{Q(s_i, a)\} .$$

- **Pro:** We do not need a model or a way to predict  $s_{i+1}$ .
- **Con:** We need to learn a  $Q$ -factor for every state/action pair.

We can learn  $Q$ -factors using a TD approach given a trajectory  $s_0, a_0, r_0, s_1, a_1, r_1 \dots, s_T, r_T$ :

$$\hat{Q}(s_i, a_i) = r_i + \gamma Q(s_{i+1}, a_{i+1}) \quad \text{target}$$

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha \left[ \hat{Q}(s_i, a_i) - Q(s_i, a_i) \right] \quad \text{update}$$

This approach is called *SARSA*.

## SARSA follows a trajectory, not an optimal path

The SARSA update equation is

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha \left[ \underbrace{r_i + \gamma Q(s_{i+1}, a_{i+1})}_{\text{target}} - Q(s_i, a_i) \right].$$

## SARSA follows a trajectory, not an optimal path

The SARSA update equation is

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha \left[ \underbrace{r_i + \gamma Q(s_{i+1}, a_{i+1})}_{\text{target}} - Q(s_i, a_i) \right].$$

Our estimate of  $Q(s_i, a_i)$  is based on

- ▶ The reward  $r_i$  experienced by selecting action  $a_i$  in state  $s_i$ .
- ▶ The future reward  $Q(s_{i+1}, a_{i+1})$  based on the action  $a_{i+1}$  from the trajectory.



## SARSA follows a trajectory, not an optimal path

The SARSA update equation is

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha \left[ \underbrace{r_i + \gamma Q(s_{i+1}, a_{i+1})}_{\text{target}} - Q(s_i, a_i) \right].$$

Our estimate of  $Q(s_i, a_i)$  is based on

- ▶ The reward  $r_i$  experienced by selecting action  $a_i$  in state  $s_i$ .
- ▶ The future reward  $Q(s_{i+1}, a_{i+1})$  based on the action  $a_{i+1}$  from the trajectory.

The policy that generated the trajectory is not optimal, so it is likely that  $a_{i+1}$  was not the best action to take.

Selecting a suboptimal action underestimates the reward to go, and therefore the value  $Q(s_i, a_i)$ .

## Q-learning

The Q-learning algorithm changes the SARSA update

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha [r_i + \gamma Q(s_{i+1}, a_{i+1}) - Q(s_i, a_i)]$$

to use the optimal action in state  $s_{i+1}$ :

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha \left[ r_i + \gamma \max_a Q(s_{i+1}, a) - Q(s_i, a_i) \right].$$

## Q-learning

The Q-learning algorithm changes the SARSA update

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha [r_i + \gamma Q(s_{i+1}, a_{i+1}) - Q(s_i, a_i)]$$

to use the optimal action in state  $s_{i+1}$ :

$$Q(s_i, a_i) \leftarrow Q(s_i, a_i) + \alpha \left[ r_i + \gamma \max_a Q(s_{i+1}, a) - Q(s_i, a_i) \right].$$

Q-learning can converge faster to an optimal policy. However, it has two drawbacks:

1. If the number of available actions is large, the maximization operator can be expensive to evaluate.
2. The maximization operator is biased.

Records were meant to be broken.

- ▶ Imagine that the quality of professional basketball players was fixed over time.
- ▶ In this case, scoring records would still be broken.
- ▶ Basketball includes stochastic elements, so as more games are played the chance of observing outliers increases.

Records were meant to be broken.

- ▶ Imagine that the quality of professional basketball players was fixed over time.
- ▶ In this case, scoring records would still be broken.
- ▶ Basketball includes stochastic elements, so as more games are played the chance of observing outliers increases.

Any algorithm with a `max` operator will drift upwards over time, *even if the mean value remains fixed*.

For  $Q$ -learning, we need to combat the bias in the `max` operator.

## Double $Q$ -learning

One solution to the max bias is using two separate  $Q$  functions (networks), called  $Q_1$  and  $Q_2$ .

Both  $Q_1$  and  $Q_2$  are trained with separate experiences. (Or, one network can *lag* behind the other in experiences.)

## Double Q-learning

One solution to the max bias is using two separate  $Q$  functions (networks), called  $Q_1$  and  $Q_2$ .

Both  $Q_1$  and  $Q_2$  are trained with separate experiences. (Or, one network can lag behind the other in experiences.)

When updating, we use one network to select the action, and the other network to compute its value.

$$\begin{aligned}Q_1(s_i, a_i) &\leftarrow Q_1(s_i, a_i) + \alpha [r_i + \gamma Q_2(s_{i+1}, a_1) - Q_1(s_i, a_i)] \\a_1 &\equiv \arg \max_a Q_1(s_{i+1}, a)\end{aligned}$$

$$\begin{aligned}Q_2(s_i, a_i) &\leftarrow Q_2(s_i, a_i) + \alpha [r_i + \gamma Q_1(s_{i+1}, a_2) - Q_2(s_i, a_i)] \\a_2 &\equiv \arg \max_a Q_2(s_{i+1}, a)\end{aligned}$$

## Double Q-learning

One solution to the max bias is using two separate  $Q$  functions (networks), called  $Q_1$  and  $Q_2$ .

Both  $Q_1$  and  $Q_2$  are trained with separate experiences. (Or, one network can lag behind the other in experiences.)

When updating, we use one network to select the action, and the other network to compute its value.

$$\begin{aligned}Q_1(s_i, a_i) &\leftarrow Q_1(s_i, a_i) + \alpha [r_i + \gamma Q_2(s_{i+1}, a_1) - Q_1(s_i, a_i)] \\a_1 &\equiv \arg \max_a Q_1(s_{i+1}, a)\end{aligned}$$

$$\begin{aligned}Q_2(s_i, a_i) &\leftarrow Q_2(s_i, a_i) + \alpha [r_i + \gamma Q_1(s_{i+1}, a_2) - Q_2(s_i, a_i)] \\a_2 &\equiv \arg \max_a Q_2(s_{i+1}, a)\end{aligned}$$

Even if  $a_1$  was selected because  $Q_1(s_{i+1}, a_1)$  was aberrantly high, the value  $Q_2(s_{i+1}, a_1)$  will not share this bias.



## Deep $Q$ -learning

- ▶ Currently, the most common method for approximating  $Q$ -factors is *deep learning* with artificial neural networks.
- ▶ We're going to learn to play a simple board game called Tic-Tac-Go.
- ▶ We want a game that is simple enough to be computationally tractable, but not easily solved.

## Deep $Q$ -learning

- ▶ Currently, the most common method for approximating  $Q$ -factors is *deep learning* with artificial neural networks.
- ▶ We're going to learn to play a simple board game called Tic-Tac-Go.
- ▶ We want a game that is simple enough to be computationally tractable, but not easily solved.
- ▶ Tic-Tac-Toe is simple, but solved. If both players follow an optimal strategy, the game will always end in a draw.

## Deep $Q$ -learning

- ▶ Currently, the most common method for approximating  $Q$ -factors is *deep learning* with artificial neural networks.
- ▶ We're going to learn to play a simple board game called Tic-Tac-Go.
- ▶ We want a game that is simple enough to be computationally tractable, but not easily solved.
- ▶ Tic-Tac-Toe is simple, but solved. If both players follow an optimal strategy, the game will always end in a draw.
- ▶ Go is unsolved, but approximating  $Q$ -factors is ridiculously expensive.


# Tic-Tac-Go

- ▶ Tic-Tac-Go is played on a  $4 \times 4$  grid.
- ▶ Before playing, 4 squares are randomly “blocked”.
- ▶ Two players, X and O, alternate placing pieces in open squares.
- ▶ Players receive points for making horizontal or vertical “chains”.
- ▶ A chain of length  $k$  is worth  $(k - 1)^2$  points.

length 2 = 1 point

length 3 = 4 points

length 4 = 9 points



# Tic-Tac-Go

- ▶ Tic-Tac-Go is played on a  $4 \times 4$  grid.
- ▶ Before playing, 4 squares are randomly “blocked”.
- ▶ Two players, X and O, alternate placing pieces in open squares.
- ▶ Players receive points for making horizontal or vertical “chains”.
- ▶ A chain of length  $k$  is worth  $(k - 1)^2$  points.

length 2 = 1 point

length 3 = 4 points

length 4 = 9 points

O	O	X	O
	O	X	X
	X	X	
O		O	X

# Tic-Tac-Go

- ▶ Tic-Tac-Go is played on a  $4 \times 4$  grid.
- ▶ Before playing, 4 squares are randomly “blocked”.
- ▶ Two players, X and O, alternate placing pieces in open squares.
- ▶ Players receive points for making horizontal or vertical “chains”.
- ▶ A chain of length  $k$  is worth  $(k - 1)^2$  points.

length 2 = 1 point

length 3 = 4 points

length 4 = 9 points

O	O	X	O
	O	X	X
	X	X	
O		O	X

X score:

$$(2-1)^2 + (2-1)^2 + (3-1)^2 = 6$$

O score:

$$(2-1)^2 + (2-1)^2 = 2$$

## States for Tic-Tac-Go

- ▶ States  $s_i$  are configurations of the board.
- ▶ Each of the 16 squares can be empty, blocked, X, or O.
- ▶ There are  ${}_{16}C_4 \times 3^{12} = 967,222,620$  possible board configurations.
- ▶ Each configuration has, on average, 6 possible moves, so there are more than  $5.8 \times 10^9$   $Q$ -factors to learn!

## States for Tic-Tac-Go

- ▶ States  $s_i$  are configurations of the board.
- ▶ Each of the 16 squares can be empty, blocked, X, or O.
- ▶ There are  ${}_{16}C_4 \times 3^{12} = 967,222,620$  possible board configurations.
- ▶ Each configuration has, on average, 6 possible moves, so there are more than  $5.8 \times 10^9$   $Q$ -factors to learn!

How do we encode the states for a function approximator?



## Option 1: Ignore the blocked states, $-1$ , $0$ , $+1$

- ▶ Let's ignore the blocked squares since we can't play on them.
- ▶ A square with X is  $-1$ , 0 is  $+1$ , and empty squares are  $0$ .
- ▶ Each state  $s_i$  is a  $12 \times 1$  trinary vectory.

## Option 1: Ignore the blocked states, $-1, 0, +1$

- ▶ Let's ignore the blocked squares since we can't play on them.
- ▶ A square with X is  $-1$ , 0 is  $+1$ , and empty squares are 0.
- ▶ Each state  $s_i$  is a  $12 \times 1$  trinary vectory.

O	O		
	O		X
	X		
			X

	O		
O	O		X
	X		
			X

## Option 2: One-hot encoding

state

	O		
O	O		X
	X		
			X

blocked

1			
1			1
	1		

player 0

	1		
1	1		

empty

		1	1
		1	
		1	
1		1	

player X

			1
	1		
			1

## Option 2: One-hot encoding

state

	O		
O	O		X
	X		
			X

Each state is a  
 $16 \times 4$  matrix  
or a  
 $64 \times 1$  vector.

blocked

1			
1			1
	1		

player 0

	1		
1	1		

empty

		1	1
		1	
		1	
1		1	

player X

			1
	1		
			1

## Our plan for Tic-Tac-Go

- ▶ We'll start with a one-hot encoding, as it simplifies the neural network.
- ▶ The four features are not independent, but we leave them to accelerate learning.
- ▶ However, *state unrolling* loses spatial information about the board. We will eventually avoid unrolling via *convolution*.
- ▶ Convolution will also exploit symmetry in the game board.

## Summary

- ▶  $Q$ -learning is a state-of-the-art technique for RL.
- ▶ Double  $Q$ -learning counteracts the bias in the  $\max$  operator.
- ▶ Defining the state space for simple board games is not trivial. Some state space representations are better for learning.

## Summary

- ▶  $Q$ -learning is a state-of-the-art technique for RL.
- ▶ Double  $Q$ -learning counteracts the bias in the  $\max$  operator.
- ▶ Defining the state space for simple board games is not trivial. Some state space representations are better for learning.
- ▶ **Next time:** Artificial neural networks.